

# 컴퓨터비전-Object Detection 기말과제

## 목차

- 컴퓨터비전-Object Detection 기말과제
  - 라이브러리 import 및 설정
  - 데이터 탐색
  - YOLO 형식으로 변환
  - YOLO11n 모델 로드 및 학습
  - Model Evaluation
  - Visualization and Testing

## 라이브러리 import 및 설정

```
In [ ]: import os
import shutil
import random
import numpy as np
import matplotlib.pyplot as plt
from PIL import Image
from tqdm import tqdm
from sklearn.model_selection import train_test_split
import torch
import cv2
import yaml

# YOLO 라이브러리
from ultralytics import YOLO
import warnings

warnings.filterwarnings("ignore")
```

```

# 재현 가능한 결과를 위한 시드 설정
random.seed(42)
np.random.seed(42)
torch.manual_seed(42)

# 디바이스 설정 (Apple Silicon MPS)
if torch.backends.mps.is_available():
    device = "mps"
    print("MPS (Apple Silicon) device available")
elif torch.cuda.is_available():
    device = "cuda"
    print("CUDA device available")
else:
    device = "cpu"
    print("Using CPU device")

print(f"Using device: {device}")

```

✓ MPS (Apple Silicon) device available  
Using device: mps

## 데이터 탐색

```

In [2]: # Check data structure
image_dir = "Images"
label_dir = "Labels"

if not os.path.exists(image_dir) or not os.path.exists(label_dir):
    print("Images or Labels directory not found!")
    print("Please ensure 'Images' and 'Labels' directories exist in the current path.")
else:
    image_files = [f for f in os.listdir(image_dir) if f.endswith(".jpeg")]
    label_files = [f for f in os.listdir(label_dir) if f.endswith(".txt")]

    print(f"Found {len(image_files)} images and {len(label_files)} labels")

# Check sample annotation format
if label_files:
    sample_label = os.path.join(label_dir, label_files[0])
    with open(sample_label, "r") as f:

```

```

        content = f.read().strip()
        print(f"\nSample annotation ({label_files[0]}):")
        print(content)

    # Check sample image
    if image_files:
        sample_img = Image.open(os.path.join(image_dir, image_files[0]))
        print(f"\nSample image size: {sample_img.size}")

```

Found 333 images and 333 labels

Sample annotation (289.txt):

```

1
33 58 118 113

```

Sample image size: (243, 207)

## YOLO 형식으로 변환

```

In [3]: def convert_bbox_to_yolo(img_width, img_height, x_min, y_min, x_max, y_max):
        """
        Convert bounding box from absolute coordinates to YOLO format
        YOLO format: class x_center y_center width height (all normalized 0-1)
        """

        # Calculate center coordinates
        x_center = (x_min + x_max) / 2.0
        y_center = (y_min + y_max) / 2.0

        # Calculate width and height
        width = x_max - x_min
        height = y_max - y_min

        # Normalize to 0-1 range
        x_center_norm = x_center / img_width
        y_center_norm = y_center / img_height
        width_norm = width / img_width
        height_norm = height / img_height

        return x_center_norm, y_center_norm, width_norm, height_norm

```

```

def create_yolo_dataset():
    """
    Create YOLO format dataset from custom annotation format
    """
    # Create output directory
    yolo_dir = "gun_dataset_yolo"
    os.makedirs(yolo_dir, exist_ok=True)

    # Get all valid image-label pairs
    image_files = [f for f in os.listdir(image_dir) if f.endswith(".jpeg")]
    valid_pairs = []

    for img_file in image_files:
        img_id = os.path.splitext(img_file)[0]
        label_file = f"{img_id}.txt"

        if os.path.exists(os.path.join(label_dir, label_file)):
            valid_pairs.append((img_file, label_file))

    print(f"Processing {len(valid_pairs)} valid image-label pairs...")

    # Split data: 70% train, 15% val, 15% test
    train_pairs, temp_pairs = train_test_split(
        valid_pairs, test_size=0.3, random_state=42
    )
    val_pairs, test_pairs = train_test_split(temp_pairs, test_size=0.5, random_state=42)

    print(f"Dataset split:")
    print(
        f"    Training: {len(train_pairs)} images ({len(train_pairs)/len(valid_pairs)*100:.1f}%"
    )
    print(
        f"    Validation: {len(val_pairs)} images ({len(val_pairs)/len(valid_pairs)*100:.1f}%"
    )
    print(
        f"    Test: {len(test_pairs)} images ({len(test_pairs)/len(valid_pairs)*100:.1f}%"
    )

    # Create directories for each split
    splits = {"train": train_pairs, "val": val_pairs, "test": test_pairs}

```

[illegible]

```

        # Validate coordinates
        if (
            x_max > x_min
            and y_max > y_min
            and x_min >= 0
            and y_min >= 0
            and x_max <= img_width
            and y_max <= img_height
        ):

            # Convert to YOLO format
            x_center, y_center, width, height = (
                convert_bbox_to_yolo(
                    img_width,
                    img_height,
                    x_min,
                    y_min,
                    x_max,
                    y_max,
                )
            )

            # Class 0 for gun
            yolo_line = f"0 {x_center:.6f} {y_center:.6f} {width:.6f} {height:.6f}"
            yolo_annotations.append(yolo_line)

    except ValueError:
        print(f"⚠ Warning: Invalid format in {label_file}")

    # Save YOLO format annotation
    with open(dst_label, "w") as f:
        f.write("\n".join(yolo_annotations))

# Create data.yaml file
data_yaml = {
    "path": yolo_dir,
    "train": "train/images",
    "val": "val/images",
    "test": "test/images",
    "names": {0: "gun"},

```

```

        "nc": 1,
    }

    with open(os.path.join(yolo_dir, "data.yaml"), "w") as f:
        yaml.dump(data_yaml, f, default_flow_style=False)

    print(f"\nYOLO dataset created successfully!")
    print(f"Dataset location: {yolo_dir}/")
    print(f"Configuration file: {yolo_dir}/data.yaml")

    return yolo_dir

# Create the YOLO dataset
dataset_path = create_yolo_dataset()

```

Processing 333 valid image-label pairs...

Dataset split:

Training: 233 images (70.0%)

Validation: 50 images (15.0%)

Test: 50 images (15.0%)

Converting train set...

Converting train: 100%|██████████| 233/233 [00:00<00:00, 1553.96it/s]

Converting val set...

Converting val: 100%|██████████| 50/50 [00:00<00:00, 1715.50it/s]

Converting test set...

Converting test: 100%|██████████| 50/50 [00:00<00:00, 1293.13it/s]

YOLO dataset created successfully!

Dataset location: gun\_dataset\_yolo/

Configuration file: gun\_dataset\_yolo/data.yaml

## YOLO11n 모델 로드 및 학습

```

In [4]: # Load YOLO11n model
        model = YOLO("yolo11s.pt")

        print("\nModel Information:")

```

```
print(f"Model: {model.model_name if hasattr(model, 'model_name') else 'YOLO11s'}")
print(f"Task: Object Detection")
print(f"Classes: 1 (gun)")
```

Model Information:

Model: yolo11s.pt

Task: Object Detection

Classes: 1 (gun)

```
In [ ]: # Train the model
# Training configuration
results = model.train(
    data=f"{dataset_path}/data.yaml", # Path to dataset configuration
    epochs=50, # Number of epochs
    imgsz=180, # Image size
    batch=8, # Batch size (adjust based on your memory)
    device=device, # Use MPS/CUDA/CPU
    project="runs/detect", # Output directory
    name="gun_detection_yolo11s_optimized", # Experiment name
    save=True, # Save checkpoints
    patience=15, # Early stopping patience (increased)
    optimizer="AdamW", # Optimizer
    lr0=0.001, # Initial learning rate
    weight_decay=0.0005, # Weight decay
    warmup_epochs=3, # Warmup epochs
    val=True, # Enable validation
    plots=True, # Generate training plots
    seed=42, # Random seed
    deterministic=True, # Deterministic training
    # NMS and performance optimization
    iou=0.6, # NMS IoU threshold (reduced for faster NMS)
    max_det=100, # Maximum detections per image (reduced)
    workers=2, # Data loading workers
    rect=True, # Rectangular training for efficiency
)

print(f"Results saved to: {results.save_dir}")

# Get the best model path
best_model_path = os.path.join(results.save_dir, "weights", "best.pt")
print(f"Best model saved at: {best_model_path}")
```



Ultralytics 8.3.151 🚀 Python-3.11.6 torch-2.7.1 MPS (Apple M3 Pro)

**engine/trainer:** agnostic\_nms=False, amp=True, augment=False, auto\_augment=randaugument, batch=8, bgr=0.0, box=7.5, cache=False, cfg=None, classes=None, close\_mosaic=10, cls=0.5, conf=None, copy\_paste=0.0, copy\_paste\_mode=flip, cos\_lr=False, cutmix=0.0, data=gun\_dataset\_yolo/data.yaml, degrees=0.0, deterministic=True, device=mps, dfl=1.5, dnn=False, dropout=0.0, dynamic=False, embed=None, epochs=50, erasing=0.4, exist\_ok=False, flipplr=0.5, flipud=0.0, format=torchscript, fraction=1.0, freeze=None, half=False, hsv\_h=0.015, hsv\_s=0.7, hsv\_v=0.4, imgsz=180, int8=False, iou=0.6, keras=False, kobj=1.0, line\_width=None, lr0=0.001, lrf=0.01, mask\_ratio=4, max\_det=100, mixup=0.0, mode=train, model=yolo11s.pt, momentum=0.937, mosaic=1.0, multi\_scale=False, name=gun\_detection\_yolo11s\_optimized, nbs=64, nms=False, opset=None, optimize=False, optimizer=AdamW, overlap\_mask=True, patience=15, perspective=0.0, plots=True, pose=12.0, pretrained=True, profile=False, project=runs/detect, rect=True, resume=False, retina\_masks=False, save=True, save\_conf=False, save\_crop=False, save\_dir=runs/detect/gun\_detection\_yolo11s\_optimized, save\_frames=False, save\_json=False, save\_period=-1, save\_txt=False, scale=0.5, seed=42, shear=0.0, show=False, show\_boxes=True, show\_conf=True, show\_labels=True, simplify=True, single\_cls=False, source=None, split=val, stream\_buffer=False, task=detect, time=None, tracker=botsort.yaml, translate=0.1, val=True, verbose=True, vid\_stride=1, visualize=False, warmup\_bias\_lr=0.1, warmup\_epochs=3, warmup\_momentum=0.8, weight\_decay=0.0005, workers=2, workspace=None

Overriding model.yaml nc=80 with nc=1

	from	n	params	module	arguments
0	-1	1	928	ultralytics.nn.modules.conv.Conv	[3, 32, 3, 2]
1	-1	1	18560	ultralytics.nn.modules.conv.Conv	[32, 64, 3, 2]
2	-1	1	26080	ultralytics.nn.modules.block.C3k2	[64, 128, 1, False, 0.25]
3	-1	1	147712	ultralytics.nn.modules.conv.Conv	[128, 128, 3, 2]
4	-1	1	103360	ultralytics.nn.modules.block.C3k2	[128, 256, 1, False, 0.25]
5	-1	1	590336	ultralytics.nn.modules.conv.Conv	[256, 256, 3, 2]
6	-1	1	346112	ultralytics.nn.modules.block.C3k2	[256, 256, 1, True]
7	-1	1	1180672	ultralytics.nn.modules.conv.Conv	[256, 512, 3, 2]
8	-1	1	1380352	ultralytics.nn.modules.block.C3k2	[512, 512, 1, True]
9	-1	1	656896	ultralytics.nn.modules.block.SPPF	[512, 512, 5]
10	-1	1	990976	ultralytics.nn.modules.block.C2PSA	[512, 512, 1]
11	-1	1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
12	[-1, 6]	1	0	ultralytics.nn.modules.conv.Concat	[1]
13	-1	1	443776	ultralytics.nn.modules.block.C3k2	[768, 256, 1, False]
14	-1	1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
15	[-1, 4]	1	0	ultralytics.nn.modules.conv.Concat	[1]
16	-1	1	127680	ultralytics.nn.modules.block.C3k2	[512, 128, 1, False]
17	-1	1	147712	ultralytics.nn.modules.conv.Conv	[128, 128, 3, 2]
18	[-1, 13]	1	0	ultralytics.nn.modules.conv.Concat	[1]
19	-1	1	345472	ultralytics.nn.modules.block.C3k2	[384, 256, 1, False]
20	-1	1	590336	ultralytics.nn.modules.conv.Conv	[256, 256, 3, 2]
21	[-1, 10]	1	0	ultralytics.nn.modules.conv.Concat	[1]
22	-1	1	1511424	ultralytics.nn.modules.block.C3k2	[768, 512, 1, True]

23 [16, 19, 22] 1 819795 ultralytics.nn.modules.head.Detect [1, [128, 256, 512]]  
YOLO11s summary: 181 layers, 9,428,179 parameters, 9,428,163 gradients, 21.5 GFLOPs

Transferred 493/499 items from pretrained weights

Freezing layer 'model.23.dfl.conv.weight'

WARNING ⚠️ imgsz=[180] must be multiple of max stride 32, updating to [192]

**train:** Fast image access ✅ (ping: 0.0±0.0 ms, read: 259.0±68.2 MB/s, size: 5.5 KB)

**train:** Scanning /Users/syshin/Downloads/archive/gun\_dataset\_yolo/train/labels.cache... 233 images, 0 backgrounds, 0 corrupt: 100%|██████████| 233/233 [00:00<?, ?it/s]

WARNING ⚠️ 'rect=True' is incompatible with DataLoader shuffle, setting shuffle=False

**val:** Fast image access ✅ (ping: 0.0±0.0 ms, read: 376.6±180.6 MB/s, size: 6.5 KB)

**val:** Scanning /Users/syshin/Downloads/archive/gun\_dataset\_yolo/val/labels.cache... 50 images, 0 backgrounds, 0 corrupt: 100%|██████████| 50/50 [00:00<?, ?it/s]

Plotting labels to runs/detect/gun\_detection\_yolo11s\_optimized/labels.jpg...

**optimizer:** AdamW(lr=0.001, momentum=0.937) with parameter groups 81 weight(decay=0.0), 88 weight(decay=0.0005), 87 bias(decay=0.0)

Image sizes 192 train, 192 val

Using 0 dataloader workers

Logging results to runs/detect/gun\_detection\_yolo11s\_optimized

Starting training for 50 epochs...

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size		
1/50	0.802G	2.3	2.354	1.765	1	128: 100% ██████████  30/30 [00:25<00:00, 1.17it/s]		
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:12<00:00, 3.16s/it]	
	all	50	66	0.0904	0.136	0.0564	0.0184	
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size		
2/50	0.82G	2.211	1.98	1.692	1	128: 100% ██████████  30/30 [00:12<00:00, 2.45it/s]		
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:10<00:00, 2.57s/it]	
	all	50	66	0.212	0.197	0.166	0.0676	
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size		
3/50	0.822G	1.978	1.83	1.485	1	128: 100% ██████████  30/30 [00:11<00:00, 2.61it/s]		
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:07<00:00, 1.94s/it]	

all		50	66	0.398	0.455	0.375	0.159
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size	
4/50	0.808G	1.986	1.659	1.582	1	128:	100% ██████████  30/30 [00:11<00:00, 2.65it/s]
Class		Images	Instances	Box(P	R	mAP50	mAP50-95): 0% ██████████  0/4 [00:00<?, ?it/s]
WARNING ⚠️ NMS time limit 2.800s exceeded							
Class		Images	Instances	Box(P	R	mAP50	mAP50-95): 25% ██████████  1/4 [00:05<00:16, 5.61s/it]
WARNING ⚠️ NMS time limit 2.800s exceeded							
Class		Images	Instances	Box(P	R	mAP50	mAP50-95): 50% ██████████  2/4 [00:10<00:10, 5.22s/it]
WARNING ⚠️ NMS time limit 2.800s exceeded							
Class		Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:15<00:00, 3.93s/it]
all		50	66	0.364	0.348	0.312	0.112
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size	
5/50	0.809G	1.889	1.565	1.517	1	128:	100% ██████████  30/30 [00:12<00:00, 2.45it/s]
Class		Images	Instances	Box(P	R	mAP50	mAP50-95): 0% ██████████  0/4 [00:00<?, ?it/s]
WARNING ⚠️ NMS time limit 2.800s exceeded							
Class		Images	Instances	Box(P	R	mAP50	mAP50-95): 25% ██████████  1/4 [00:05<00:16, 5.56s/it]
WARNING ⚠️ NMS time limit 2.800s exceeded							
Class		Images	Instances	Box(P	R	mAP50	mAP50-95): 50% ██████████  2/4 [00:09<00:09, 4.74s/it]
WARNING ⚠️ NMS time limit 2.800s exceeded							
Class		Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:16<00:00, 4.24s/it]
all		50	66	0.417	0.423	0.295	0.0923
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size	
6/50	0.81G	1.93	1.556	1.458	1	128:	100% ██████████  30/30 [00:13<00:00, 2.15it/s]
Class		Images	Instances	Box(P	R	mAP50	mAP50-95): 0% ██████████  0/4 [00:00<?, ?it/s]
WARNING ⚠️ NMS time limit 2.800s exceeded							

6<00:19, 6.61s/it]	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 25%	1/4 [00:0
WARNING ⚠ NMS time limit 2.800s exceeded								
1<00:10, 5.49s/it]	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 50%	2/4 [00:1
WARNING ⚠ NMS time limit 2.800s exceeded								
7<00:00, 4.41s/it]	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100%	4/4 [00:1
all		50	66	0.621	0.439	0.503	0.202	
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size		
7/50	0.819G	1.858	1.483	1.443	1	128: 100%	30/30 [00:14<00:00,	
2.08it/s]	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 0%	0/4 [00:0
0<?, ?it/s]								
WARNING ⚠ NMS time limit 2.800s exceeded								
5<00:17, 5.74s/it]	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 25%	1/4 [00:0
WARNING ⚠ NMS time limit 2.800s exceeded								
6<00:00, 4.06s/it]	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100%	4/4 [00:1
all		50	66	0.575	0.369	0.37	0.178	
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size		
8/50	0.813G	1.749	1.371	1.356	1	128: 100%	30/30 [00:15<00:00,	
1.99it/s]	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 0%	0/4 [00:0
0<?, ?it/s]								
WARNING ⚠ NMS time limit 2.800s exceeded								
8<00:24, 8.32s/it]	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 25%	1/4 [00:0
WARNING ⚠ NMS time limit 2.800s exceeded								
0<00:00, 5.05s/it]	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100%	4/4 [00:2
all		50	66	0.61	0.636	0.634	0.234	
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size		

9/50	0.806G	1.745	1.419	1.393	1	128: 100% ██████████  30/30 [00:17<00:00, 1.76it/s]
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 25% ██████████  1/4 [00:05<00:15, 5.16s/it]
WARNING ⚠ NMS time limit 2.800s exceeded						
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:16<00:00, 4.17s/it]
all	50	66	0.578	0.682	0.632	0.238
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size
10/50	0.814G	1.768	1.304	1.441	1	128: 100% ██████████  30/30 [00:14<00:00, 2.05it/s]
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:11<00:00, 2.79s/it]
all	50	66	0.67	0.621	0.625	0.264
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size
11/50	0.815G	1.711	1.261	1.382	1	128: 100% ██████████  30/30 [00:16<00:00, 1.86it/s]
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:13<00:00, 3.42s/it]
all	50	66	0.693	0.606	0.628	0.236
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size
12/50	0.815G	1.645	1.262	1.377	1	128: 100% ██████████  30/30 [00:19<00:00, 1.57it/s]
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:11<00:00, 2.86s/it]
all	50	66	0.669	0.652	0.677	0.291
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size
13/50	0.808G	1.593	1.267	1.354	1	128: 100% ██████████  30/30 [00:22<00:00, 1.35it/s]
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 25% ██████████  1/4 [00:05<00:17, 5.85s/it]
WARNING ⚠ NMS time limit 2.800s exceeded						
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:18<00:00, 4.66s/it]
all	50	66	0.594	0.697	0.675	0.278
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size

14/50	0.816G	1.663	1.19	1.35	1	128: 100% ██████████  30/30 [00:21<00:00, 1.42it/s]
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:12<00:00, 3.03s/it]
all	50	66	0.592	0.682	0.686	0.262
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
15/50	0.817G	1.623	1.1	1.347	1	128: 100% ██████████  30/30 [00:15<00:00, 1.94it/s]
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:11<00:00, 2.97s/it]
all	50	66	0.745	0.576	0.696	0.286
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
16/50	0.817G	1.609	1.071	1.323	1	128: 100% ██████████  30/30 [00:17<00:00, 1.74it/s]
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:07<00:00, 1.89s/it]
all	50	66	0.734	0.652	0.709	0.285
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
17/50	0.809G	1.607	1.101	1.286	1	128: 100% ██████████  30/30 [00:18<00:00, 1.64it/s]
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:08<00:00, 2.11s/it]
all	50	66	0.668	0.606	0.621	0.237
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
18/50	0.817G	1.612	1.04	1.3	1	128: 100% ██████████  30/30 [00:29<00:00, 1.02it/s]
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:12<00:00, 3.15s/it]
all	50	66	0.729	0.682	0.714	0.294
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
19/50	0.818G	1.544	1.13	1.249	1	128: 100% ██████████  30/30 [00:23<00:00, 1.27it/s]
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:09<00:00, 2.30s/it]
all	50	66	0.884	0.697	0.787	0.335
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size

20/50 1.02s/it]	0.818G	1.503	1.021	1.283	1	128: 100% ██████████  30/30 [00:30<00:00,
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:0
9<00:00, 2.39s/it]						
all	50	66	0.742	0.699	0.753	0.339
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
21/50 1.27it/s]	0.81G	1.438	1.006	1.253	1	128: 100% ██████████  30/30 [00:23<00:00,
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:0
7<00:00, 1.88s/it]						
all	50	66	0.744	0.742	0.753	0.349
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
22/50 1.24it/s]	0.818G	1.435	0.8937	1.224	1	128: 100% ██████████  30/30 [00:24<00:00,
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:1
0<00:00, 2.65s/it]						
all	50	66	0.805	0.621	0.69	0.36
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
23/50 1.35s/it]	0.818G	1.512	0.9733	1.228	1	128: 100% ██████████  30/30 [00:40<00:00,
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:1
1<00:00, 2.99s/it]						
all	50	66	0.892	0.561	0.674	0.313
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
24/50 1.10it/s]	0.818G	1.438	0.9247	1.199	1	128: 100% ██████████  30/30 [00:27<00:00,
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:0
7<00:00, 1.98s/it]						
all	50	66	0.738	0.727	0.76	0.35
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
25/50 1.78it/s]	0.81G	1.418	0.9152	1.212	1	128: 100% ██████████  30/30 [00:16<00:00,
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:0
8<00:00, 2.11s/it]						
all	50	66	0.806	0.758	0.843	0.382
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size

26/50	0.818G	1.362	0.8815	1.172	1	128: 100% ██████████  30/30 [00:27<00:00, 1.08it/s]
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:07<00:00, 1.93s/it]
all	50	66	0.923	0.725	0.82	0.348
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size
27/50	0.818G	1.358	0.85	1.171	1	128: 100% ██████████  30/30 [00:22<00:00, 1.32it/s]
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:09<00:00, 2.39s/it]
all	50	66	0.838	0.706	0.791	0.336
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size
28/50	0.818G	1.342	0.8365	1.143	1	128: 100% ██████████  30/30 [00:22<00:00, 1.31it/s]
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:07<00:00, 1.95s/it]
all	50	66	0.812	0.712	0.79	0.332
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size
29/50	0.818G	1.311	0.8262	1.138	1	128: 100% ██████████  30/30 [00:21<00:00, 1.40it/s]
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:11<00:00, 2.97s/it]
all	50	66	0.772	0.773	0.808	0.357
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size
30/50	0.818G	1.254	0.7877	1.134	1	128: 100% ██████████  30/30 [00:24<00:00, 1.25it/s]
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:07<00:00, 1.91s/it]
all	50	66	0.843	0.712	0.78	0.366
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size
31/50	0.818G	1.233	0.8079	1.126	1	128: 100% ██████████  30/30 [00:19<00:00, 1.52it/s]
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:06<00:00, 1.64s/it]
all	50	66	0.875	0.741	0.801	0.347
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size



32/50 1.22it/s]	0.818G	1.259	0.7526	1.13	1	128: 100% ██████████  30/30 [00:24<00:00,
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:0
7<00:00, 1.77s/it]						
all	50	66	0.787	0.729	0.778	0.338
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
33/50 1.18it/s]	0.811G	1.303	0.7643	1.142	1	128: 100% ██████████  30/30 [00:25<00:00,
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:0
6<00:00, 1.61s/it]						
all	50	66	0.813	0.792	0.818	0.369
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
34/50 2.00it/s]	0.818G	1.254	0.7517	1.106	1	128: 100% ██████████  30/30 [00:14<00:00,
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:0
6<00:00, 1.68s/it]						
all	50	66	0.869	0.727	0.838	0.37
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
35/50 1.28it/s]	0.818G	1.218	0.7388	1.131	1	128: 100% ██████████  30/30 [00:23<00:00,
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:0
5<00:00, 1.46s/it]						
all	50	66	0.826	0.727	0.839	0.366
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
36/50 1.46it/s]	0.818G	1.166	0.6864	1.091	1	128: 100% ██████████  30/30 [00:20<00:00,
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:0
5<00:00, 1.50s/it]						
all	50	66	0.785	0.848	0.823	0.362
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
37/50 1.17it/s]	0.811G	1.198	0.7171	1.111	1	128: 100% ██████████  30/30 [00:25<00:00,
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:0
4<00:00, 1.24s/it]						
all	50	66	0.879	0.833	0.835	0.366
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size

38/50 1.48it/s]	0.819G	1.143	0.6768	1.082	1	128: 100% ██████████  30/30 [00:20<00:00,
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:0
6<00:00, 1.58s/it]						
all	50	66	0.856	0.81	0.825	0.369
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size
39/50 2.11it/s]	0.819G	1.232	0.7379	1.087	1	128: 100% ██████████  30/30 [00:14<00:00,
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:0
4<00:00, 1.24s/it]						
all	50	66	0.859	0.818	0.828	0.375
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size
40/50 2.96it/s]	0.819G	1.146	0.6939	1.086	1	128: 100% ██████████  30/30 [00:10<00:00,
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:0
5<00:00, 1.36s/it]						
all	50	66	0.883	0.758	0.847	0.384
Closing dataloader mosaic						
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size
41/50 1.57it/s]	0.819G	1.094	0.6776	1.058	1	128: 100% ██████████  30/30 [00:19<00:00,
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:0
4<00:00, 1.10s/it]						
all	50	66	0.795	0.864	0.863	0.389
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size
42/50 2.30it/s]	0.819G	1.089	0.6435	1.041	1	128: 100% ██████████  30/30 [00:13<00:00,
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:0
3<00:00, 1.06it/s]						
all	50	66	0.844	0.788	0.843	0.382
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size
43/50 3.14it/s]	0.819G	1.093	0.6464	1.052	1	128: 100% ██████████  30/30 [00:09<00:00,
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:0
4<00:00, 1.05s/it]						
all	50	66	0.837	0.788	0.829	0.35

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size	
44/50 3.08it/s]	0.819G	1.055	0.6437	1.041	1	128: 100% ██████████	30/30 [00:09<00:00,
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:04<00:00, 1.20s/it]	
all	50	66	0.826	0.803	0.849	0.385	
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size	
45/50 1.79it/s]	0.811G	0.9945	0.6143	1.012	1	128: 100% ██████████	30/30 [00:16<00:00,
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:05<00:00, 1.29s/it]	
all	50	66	0.792	0.807	0.836	0.392	
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size	
46/50 1.75it/s]	0.819G	1.045	0.6004	1.019	1	128: 100% ██████████	30/30 [00:17<00:00,
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:04<00:00, 1.08s/it]	
all	50	66	0.892	0.754	0.862	0.412	
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size	
47/50 3.21it/s]	0.819G	1.101	0.6171	1.051	1	128: 100% ██████████	30/30 [00:09<00:00,
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:04<00:00, 1.16s/it]	
all	50	66	0.893	0.758	0.87	0.4	
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size	
48/50 1.90it/s]	0.819G	0.9786	0.5928	1.003	1	128: 100% ██████████	30/30 [00:15<00:00,
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:04<00:00, 1.02s/it]	
all	50	66	0.897	0.796	0.866	0.39	
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size	
49/50 2.29it/s]	0.811G	1.067	0.6005	1.025	1	128: 100% ██████████	30/30 [00:13<00:00,
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  4/4 [00:03<00:00, 1.07it/s]	
all	50	66	0.881	0.782	0.853	0.39	
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size	

Class	Images	Instances	Box(P	R	mAP50	mAP50-95)
all	50	66	0.884	0.807	0.869	0.394

50 epochs completed in 0.401 hours.  
Optimizer stripped from runs/detect/gun\_detection\_yolo11s\_optimized/weights/last.pt, 19.1MB  
Optimizer stripped from runs/detect/gun\_detection\_yolo11s\_optimized/weights/best.pt, 19.1MB

Validating runs/detect/gun\_detection\_yolo11s\_optimized/weights/best.pt...  
Ultralytics 8.3.151 🚀 Python-3.11.6 torch-2.7.1 MPS (Apple M3 Pro)  
YOLO11s summary (fused): 100 layers, 9,413,187 parameters, 0 gradients, 21.3 GFLOPs

Class	Images	Instances	Box(P	R	mAP50	mAP50-95)
all	50	66	0.892	0.754	0.862	0.413

Speed: 1.0ms preprocess, 321.7ms inference, 0.0ms loss, 17.9ms postprocess per image  
Results saved to **runs/detect/gun\_detection\_yolo11s\_optimized**  
Speed: 1.0ms preprocess, 321.7ms inference, 0.0ms loss, 17.9ms postprocess per image  
Results saved to **runs/detect/gun\_detection\_yolo11s\_optimized**  
Results saved to: runs/detect/gun\_detection\_yolo11s\_optimized  
Best model saved at: runs/detect/gun\_detection\_yolo11s\_optimized/weights/best.pt

## Model Evaluation

```
In [ ]: # Load the best trained model
best_model = YOLO(best_model_path)

# Evaluate on test set
metrics = best_model.val(data=f"{dataset_path}/data.yaml", split="test", device=device)

# Print evaluation metrics
print("\nEvaluation Results:")
print(f"mAP@0.5: {metrics.box.map50:.4f}")
print(f"mAP@0.5-0.95: {metrics.box.map:.4f}")
print(f"Precision: {metrics.box.mp:.4f}")
print(f"Recall: {metrics.box.mr:.4f}")
```

Ultralytics 8.3.151 🚀 Python-3.11.6 torch-2.7.1 MPS (Apple M3 Pro)  
YOLO11s summary (fused): 100 layers, 9,413,187 parameters, 0 gradients, 21.3 GFLOPs  
**val:** Fast image access ✅ (ping: 0.1±0.1 ms, read: 61.3±70.1 MB/s, size: 6.9 KB)

**val:** Scanning /Users/syshin/Downloads/archive/gun\_dataset\_yolo/test/labels.cache... 50 images, 0 backgrounds, 0 corrupt: 100%|██████████| 50/50 [00:00<?, ?it/s]

Class	Images	Instances	Box(P	R	mAP50	mAP50-95)
8<00:00, 4.56s/it]						

all	50	57	0.931	0.772	0.848	0.384
-----	----	----	-------	-------	-------	-------

Speed: 2.1ms preprocess, 195.1ms inference, 0.0ms loss, 47.4ms postprocess per image

Results saved to /Users/syshin/Downloads/archive/runs/detect/val2

Evaluation Results:

mAP@0.5: 0.8476

mAP@0.5-0.95: 0.3836

Precision: 0.9306

Recall: 0.7719

Speed: 2.1ms preprocess, 195.1ms inference, 0.0ms loss, 47.4ms postprocess per image

Results saved to /Users/syshin/Downloads/archive/runs/detect/val2

Evaluation Results:

mAP@0.5: 0.8476

mAP@0.5-0.95: 0.3836

Precision: 0.9306

Recall: 0.7719

## Visualization and Testing

```
In [7]: # Test on sample images
test_images_dir = f"{dataset_path}/test/images"
test_images = [f for f in os.listdir(test_images_dir) if f.endswith(".jpeg")][:6]

# Create visualization
fig, axes = plt.subplots(3, 2, figsize=(15, 18))
axes = axes.ravel()

for idx, img_file in enumerate(test_images):
    if idx >= 6: # Limit to 6 images
        break
```

```

img_path = os.path.join(test_images_dir, img_file)

# Run inference
results = best_model(img_path, device=device, verbose=False)

# Load and display image
img = cv2.imread(img_path)
img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

axes[idx].imshow(img)
axes[idx].set_title(f"Detection Result: {img_file}", fontsize=12)
axes[idx].axis("off")

# Draw bounding boxes
if len(results) > 0 and results[0].boxes is not None:
    boxes = results[0].boxes
    detection_count = 0

    for i in range(len(boxes)):
        # Get box coordinates and confidence
        x1, y1, x2, y2 = boxes.xyxy[i].cpu().numpy()
        conf = boxes.conf[i].cpu().numpy()

        # Only show detections with confidence > 0.3
        if conf > 0.3:
            detection_count += 1

            # Draw rectangle
            rect = plt.Rectangle(
                (x1, y1),
                x2 - x1,
                y2 - y1,
                linewidth=2,
                edgecolor="red",
                facecolor="none",
            )
            axes[idx].add_patch(rect)

            # Add confidence label
            axes[idx].text(
                x1,

```

```
        y1 - 10,  
        f"Gun: {conf:.2f}",  
        bbox=dict(boxstyle="round", facecolor="red", alpha=0.7),  
        fontsize=10,  
        color="white",  
    )  
  
    axes[idx].set_title(f"{img_file}\nDetections: {detection_count}", fontsize=10)  
else:  
    axes[idx].set_title(f"{img_file}\nNo detections", fontsize=10)  
  
plt.tight_layout()  
plt.savefig("gun_detection_results.png", dpi=300, bbox_inches="tight")  
plt.show()  
  
print("\nVisualization saved as 'gun_detection_results.png'")
```

298.jpeg  
Detections: 1



220.jpeg  
Detections: 2



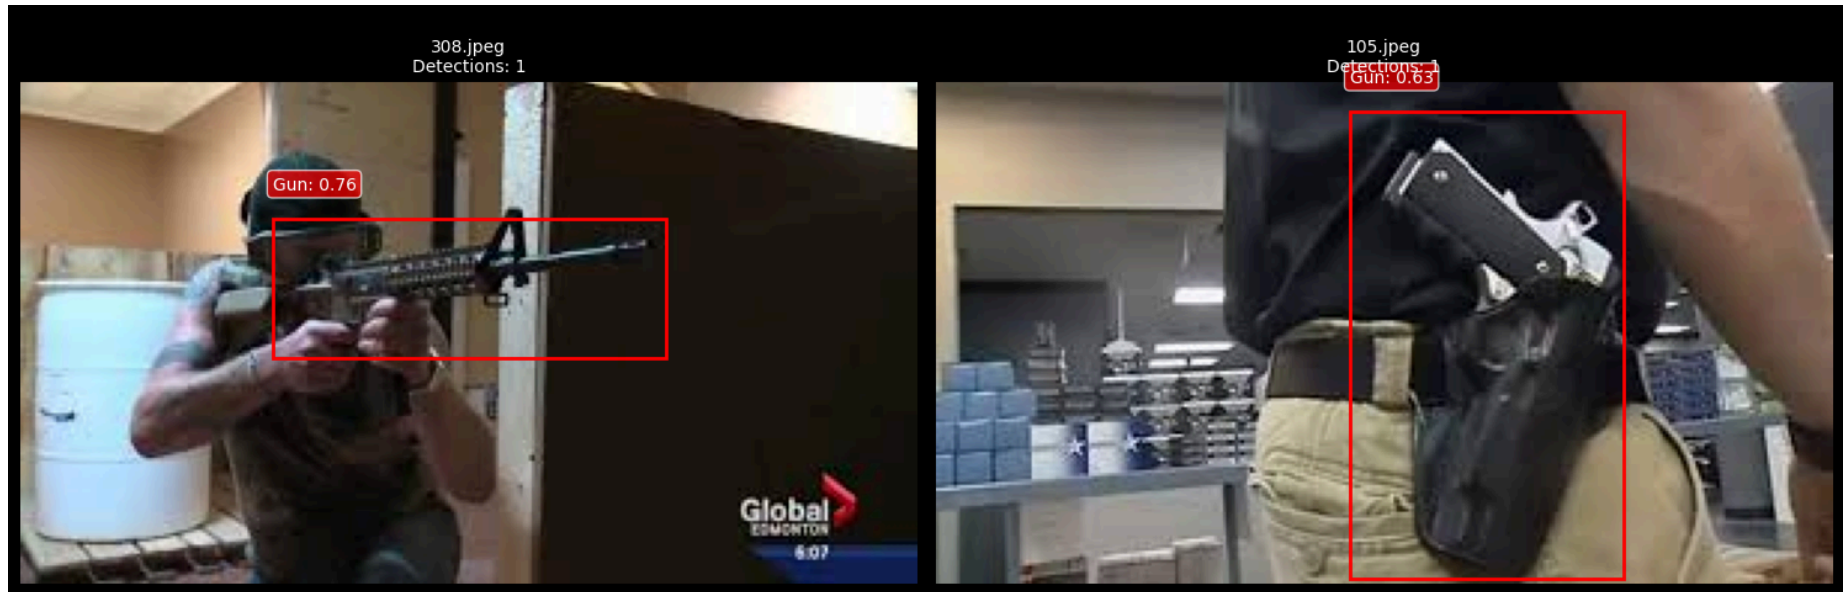
84.jpeg  
Detections: 0



133.jpeg  
Detections: 1







Visualization saved as 'gun\_detection\_results.png'

In [8]: # 최종 mAP@0.5 평가

```
print(f"mAP@0.5: {metrics.box.map50:.4f}")
```

mAP@0.5: 0.8476

### Table of contents

- 컴퓨터비전-Object Detection 기말과제
  - 라이브러리 import 및 설정
  - 데이터 탐색
  - YOLO 형식으로 변환
  - YOLO11n 모델 로드 및 학습
  - Model Evaluation
  - Visualization and Testing

In [ ]: